

Changli Han

☎ +49 15223091456 | ✉ changlihan4595@gmail.com | 📍 Berlin, Germany | [in](#)

Summary

Senior machine learning and software engineer with 8+ years building production ML products, data platforms, model serving systems, and self-service automation for analysts, applied scientists, and operations users. Strong end-to-end experience across **Python**, **SQL**, **Kubernetes**, **AWS**, CI/CD, observability, and production ownership, including model deployment, online serving, tabular ML, notebook automation, RAG/LLM APIs, and workflow platforms. Known for turning ambiguous business and user needs into scalable, reliable services with measurable impact.

Experience

DoorDash

Berlin, Germany

Machine Learning Platform Engineer

Apr. 2025 – Present

- Build and operate critical ML platform components in **Go** for feature retrieval, model deployment, and online serving used by training and inference workloads
- Develop reusable agent skills and toolsets with **Pydantic AI** for ML feature lifecycle workflows, validating tool behavior and output quality through DoorDash's internal evaluation service using built-in **Pydantic Evals** evaluators
- Own production **Kubernetes** deployments and release workflows with **GitHub Actions**, **Buildkite**, **Argo Workflows**, and **Argo CD**
- Monitor deployed services with performance metrics, alerting, and dashboards across **Datadog**, **Chronosphere**, **Grafana**, **Prometheus**, and **Rootly**
- Partner with ML and software engineers to improve self-service serving infrastructure, model deployment throughput, and production reliability

Zalando SE

Berlin, Germany

Senior Platform Engineer

Jun. 2022 – Apr. 2025

- Helped define and execute the roadmap for a compliant, cost-efficient ML experimentation platform for applied scientists and researchers
- Built **HPC** and distributed training workflows on **Kubernetes** and **AWS** with **Ray**, supporting heterogeneous CPU/GPU workloads on **NVIDIA H100/H200** environments
- Built and deployed an internal chatbot and LLM API on Kubernetes for RAG workflows, including prompt iteration support and LiteLLM-based visibility into token usage and cost governance
- Improved production observability with **Grafana** and **Prometheus**, adding service dashboards, alerts, and operational monitoring used by support teams
- Built self-service JupyterHub workspaces for research users with instance-type selection, auto-termination, SSH access, and GPU-backed options
 - * Reduced tooling costs by more than 30% through cloud cost improvements
 - * Enabled AWS SageMaker Studio for machine learning notebook users
 - * Built a **Run:ai** proof of concept for shared GPU workloads and user enablement
 - * Operated on-prem **SLURM** machines and supported research users directly

Software Engineer

Jan.2019 – Jun.2022

- Built a centralized notebook automation platform that enabled hundreds of data analysts and applied scientists to share, schedule, and operationalize analytical workflows
 - * Developed the UI with JavaScript/TypeScript, React, and Redux
 - * Built workflow orchestration with AWS Lambda, Step Functions, and EventBridge
 - * Built the Jupyter notebook extension with **Python** and Papermill
 - * Led the design and development of internal APIs in **Python** (FastAPI) backed by PostgreSQL (SQLAlchemy, Alembic)
- Supported production ML pipelines by building internal machine learning infrastructure for data scientists and applied scientists
 - * Built the front end with TypeScript, React, and Backstage

- * Developed microservices and APIs with Python and AWS SageMaker
- Built the managed internal JupyterHub SaaS used by **300+** daily active analysts and data scientists for data exploration, modeling, and business automation
 - * Provided active consulting and user support to the internal community
 - * Connected single-user notebook servers with multiple internal company data sources
 - * Automated Docker image changelogs across thousands of packages from Apt, Pip, Conda, Npm, and R sources using Python and Bash
 - * Built all-in-one Docker images providing Python and R environments with hundreds of packages
 - * Deployed JupyterHub onto the internal **Kubernetes** cluster
- Built an end-to-end tabular machine learning pipeline with a data analyst to predict Zalando Plus paid adoption and support activation and marketing decisions during the trial period
 - * Used PySpark and SQL in Databricks for data processing and feature preparation
 - * Used AWS SageMaker for model training, hyperparameter tuning, deployment, and inference
- Mentored junior engineers, analysts, and interns on software engineering, **Python**, and machine learning fundamentals.

Robert Bosch GmbH

Stuttgart, Germany

(Internship) Deep Learning Engineer - Advanced Driver-Assistance Systems

Sep.2017 – Feb.2018

- Designed and trained deep Convolutional Neural Network (CNN) models to identify drivable areas from road images using Python and Keras
- Created CNN models with average accuracy and intersection-over-union above 99% at pixel level
- Set up MkDocs documentation for the department to make project knowledge easier to maintain

Leadership and Awards

Zalando ZMS Cup

Nov.2020 – Nov.2020

- Led a group of 3 engineers and **won** the first Zalando internal Kaggle-style ML competition
- Owned data processing, feature engineering, model selection (CatBoost, LightGBM, XGBoost), and experiment tracking
- Created and presented the winning solution to the company's Data Science Community (**over 500 people**)
- Donated the prize money to a charity focused on fighting climate change

Education

RWTH Aachen University

Oct. 2015 – Apr. 2018

M.Sc in Electrical Engineering, Information Technology and Computer Engineering

Aachen, Germany

- Average Grade: 1.6

Harbin Institute of Technology

Sep.2010 – Jul.2014

Bachelor of Engineering in Communication Engineering

Harbin, China

- Average Grade: 89/100

Personal Highlights

Languages : English, Chinese, German B1

Tech : Backend: **Python** (FastAPI, pandas, scikit-learn, PySpark), **Go**; Frontend: JavaScript/TypeScript (React, Redux); Data: **SQL** (PostgreSQL, Databricks); ML/AI: model serving, model deployment, tabular ML, RAG, LLM APIs, **Pydantic AI**, **Pydantic Evals**, LiteLLM, SageMaker; MLOps/Infra: **Kubernetes**, **AWS**, **Docker**, **Terraform**, **Metaflow**, **Ray**, **SLURM**, **Run:ai**; CI/CD and GitOps: **GitHub Actions**, **Buildkite**, **Argo Workflows**, **Argo CD**; Observability/Operations: **Datadog**, **Chronosphere**, **Grafana**, **Prometheus**, **Rootly**

Operating System : **Linux**, MacOS

Certification : Functional Programming in Scala

Tools : **Git**, VSCode, Vim, Jupyter, Databricks, ...

Interests : Cloud, container, IaC, ML, CTF, clean code, ...